



Processing Scalar Implicature: A Constraint-Based Approach

Judith Degen,^a Michael K. Tanenhaus^b

^a*Department of Psychology, Stanford University*

^b*Department of Brain and Cognitive Sciences, University of Rochester*

Received 12 May 2012; received in revised form 8 July 2013; accepted 31 January 2014

Abstract

Three experiments investigated the processing of the implicature associated with *some* using a “gumball paradigm.” On each trial, participants saw an image of a gumball machine with an upper chamber with 13 gumballs and an empty lower chamber. Gumballs then dropped to the lower chamber and participants evaluated statements, such as “You got some of the gumballs.” Experiment 1 established that *some* is less natural for reference to small sets (1, 2, and 3 of the 13 gumballs) and unpartitioned sets (all 13 gumballs) compared to intermediate sets (6–8). Partitive *some of* was less natural than simple *some* when used with the unpartitioned set. In Experiment 2, including exact number descriptions lowered naturalness ratings for *some* with small sets but not for intermediate size sets and the unpartitioned set. In Experiment 3, the naturalness ratings from Experiment 2 predicted response times. The results are interpreted as evidence for a Constraint-Based account of scalar implicature processing and against both two-stage, Literal-First models and pragmatic Default models.

Keywords: Pragmatics; Scalar implicature; Quantifiers; Alternatives

1. Introduction

Successful communication requires readers and listeners (hereafter, listeners) to infer a speaker’s intended meaning from an underspecified utterance. To this end listeners make use of the semantic and pragmatic information available in the linguistic input and the discourse context. However, it is an open question whether semantic information takes a privileged position in this reasoning process. In recent years, scalar implicatures as in (1) have served as an important testing ground for examining the time course of integration of semantic and pragmatic information in inferring a speaker’s intended meaning.

Correspondence should be sent to Judith Degen, Department of Psychology, Stanford University, 450 Serra Mall, Stanford, CA 94305. E-mail: jdegen@stanford.edu

(1) You got some of the gumballs.

↪ You got some, but not all, of the gumballs

The semantic content of the utterance in (1) is that the listener got at least one gumball (lower-bound meaning of *some*). Under some conditions, however, the listener is also justified in taking the speaker to implicate that he, the listener, did not get all of the gumballs (upper-bound meaning of *some*). This is a case of *scalar implicature*. The term “scalar” is used because the inference that gives rise to the upper-bound interpretation is assumed to depend upon a highly accessible set of alternatives that are ordered on a scale by asymmetrical entailment (Horn, 2004), and that the speaker could have selected but did not (e.g., <all, some>).

1.1. Overview

The current article focuses on how listeners compute the upper-bound interpretation, sometimes called *pragmatic some*. We first discuss the classic distinction between Generalized and Particularized Conversational Implicature, originally introduced by Grice (1975) to distinguish relatively context-dependent from relatively context-independent inferences. We then describe two processing hypotheses, the Default hypothesis and the two-stage, Logical/Semantic/Literal-First hypothesis, that have played a central role in guiding experimental investigations of how listeners arrive at an interpretation of pragmatic *some*. In doing so, we flesh out some of the assumptions that underlie these approaches. We propose a probabilistic Constraint-Based framework in which naturalness and availability of alternatives (along with other factors) play a central role in computing pragmatic *some*. We then focus on two superficially similar visual world experiments that test the two-stage hypothesis, finding strikingly different data patterns and arriving at opposite conclusions. We argue that the differences arise because of the naturalness and availability of other lexical alternatives to *some* besides the stronger scalar alternative *all*, in particular the effects of intermixing *some* and exact number with small set sizes. We test this claim in three experiments using a “gumball” paradigm.

1.2. Generalized and Particularized Conversational Implicatures

According to Grice (1975), scalar implicatures are an instance of *Generalized Conversational Implicature* (GCI). These are implicatures that are assumed to arise in the same systematic way regardless of context. In the case of scalar implicature, the speaker is taken to convey the negation of a stronger alternative statement that she could have made (and that would have also been relevant) but chose not to. Instead of (1), the speaker could have said, *You got all of the gumballs*, which would have also been relevant and more informative. Together with the additional assumption that the speaker is an authority on whether or not the stronger statement is true (the Competence Assumption; Sauerland, 2004; Van Rooij & Schulz, 2004), the listener may infer

that the speaker is conveying that the listener in fact got some, but not all, of the gumballs.

Generalized Conversational Implicatures are distinguished from *Particularized Conversational Implicatures (PCI)*, which are strongly dependent on specific features of the context, such as the nature of the Question Under Discussion (QUD; Roberts, 1996). For example, if the sentence in (1) was uttered in a context in which the listener—say, a little boy—was complaining that he did not get any candy, the speaker—say, his mother—may be taken to implicate that he should stop complaining. However, this implicature does not arise if the context is modified slightly; for example, if uttered in a context where the child cannot see which objects came out of a prize machine and asks his parents whether he got all of the gumballs, the scalar implicature that he did not get all of the gumballs will presumably arise, whereas the implicature that he should stop complaining does not.

Although it is sometimes claimed that scalar implicatures arise in nearly all contexts (Levinson, 2000), scalar implicatures can, in fact, be canceled not only explicitly as in (2a) but also implicitly as in (2b).

(2a) You got some of the gumballs. In fact, you got all of them.

(2b) A: Did I get anything from the prize machine?

B: You got some of the gumballs.

In (2a), the implicature is explicitly canceled because it is immediately superseded by assertion of the stronger alternative. In (2b), it is canceled implicitly because the preceding context makes the stronger alternative with *all* irrelevant. What matters to A in (2b) is not whether he got all versus some but all of the gumballs, but whether he got any items at all (e.g., gumballs, rings, smarties) from the machine. This feature of implicit cancelability has played a crucial role in experimental studies of implicatures (e.g., Bott & Noveck, 2004; Breheny, Katsos, & Williams, 2006) because hypotheses about the time course of scalar implicature processing make different claims about the status of implicit implicature cancelation.

1.3. Processing frameworks

Formal accounts of conversational implicatures typically do not specify how a listener might compute an implicature as an utterance unfolds in a specific conversational context (e.g., Gazdar, 1979; Horn, 1984). Recently, however, there has been increased interest in how listeners compute scalar implicatures in real-time language comprehension. Most of this work has focused on distinguishing between two alternative hypotheses.

The first hypothesis is that implicatures are computed immediately and effortlessly due to their status as default inferences (Levinson, 2000; see also Chierchia, 2004 for a related account under which scalar implicatures are taken to arise by default in

upward-entailing, but not in downward-entailing contexts). Under this Default account, scalar implicatures do not incur the usual processing costs associated with generating an inference. Only cancelation of an implicature takes time and processing resources. This view is motivated by considerations about communicative efficiency. Rapid, effortless inference processes for deriving GCIs are proposed as a solution to the problem of the articulatory bottleneck: While humans can only produce a highly limited number of phonemes per second, communication nevertheless proceeds remarkably quickly.

The inferences that allow listeners to derive context-dependent interpretations, such as those presumably involved in computing particularized implicatures, are assumed to be slow and resource-intensive, in the sense of classic two-process models of attention that distinguish between *automatic* processes, which are fast, require few resources, and arise independent of context, and *controlled processes*, which are slow, strategic, and resource demanding (Neely, 1977; Posner & Snyder, 1975; Shiffrin & Schneider, 1977; also see Kahneman, 2011 for a related framework). In contrast, a scale containing a small set of lexical alternatives—like <all, some>—could be pre-compiled and thus automatically accessed, regardless of context. Because the default interpretation arises automatically, there will be a costly interpretive garden-path when the default meaning has to be overridden. Therefore, the default model predicts that in all contexts, a default, upper-bound interpretation will precede a possible lower-bound interpretation.

The second hypothesis, also termed the Literal-First hypothesis (Huang & Snedeker, 2009), assumes that the lower-bound semantic interpretation is computed rapidly, and perhaps automatically, as a by-product of basic sentence processing. All inferences, including GCIs, require extra time and resources. Therefore, the Literal-First hypothesis predicts that in all contexts a lower-bound interpretation will be computed before an upper-bound interpretation is considered. For proponents of the Literal-First hypothesis this follows from the traditional observation in the linguistic literature (e.g., Horn, 2004) that the semantic interpretation of simple declaratives containing *some* is in an important sense more basic than the pragmatic interpretation: The upper-bound interpretation *some but not all* always entails the lower-bound interpretation *at least one*.¹ The pragmatic interpretation cannot exist without the semantic one, which translates into a two-stage processing sequence: Upon encountering a scalar item like *some*, the semantic interpretation is necessarily constructed before the pragmatic one. To the extent that there is a processing distinction between Generalized and Particularized Implicatures, it is that the relevant dimensions of the context and the interpretations that drive the inference are more circumscribed and thus perhaps more accessible for Generalized Implicatures.

The Default and two-stage models make straightforward predictions about the time course of the implicature associated with the upper-bound interpretation of utterances containing *some*. The Default hypothesis predicts that logical/semantic *some* should be more resource-intensive and slower than upper-bound, pragmatic *some*, whereas the two-stage model makes the opposite prediction.

There are, however, alternative Context-Driven frameworks in which the speed with which a scalar implicature is computed is largely determined by context (e.g., Breheny

et al., 2006). For example, Relevance Theory (Sperber & Wilson, 1995), like the Literal-First hypothesis, assumes that the semantic interpretation is basic. In Relevance Theory, the upper-bound meaning is only computed if required to reach a certain threshold of relevance in context. In contrast to the Literal-First hypothesis, however, Relevance Theory does not necessarily assume a processing cost for the pragmatic inference. If the context provides sufficient support for the upper-bound interpretation, it may well be computed without incurring additional processing cost. However, a processing cost will be incurred if the upper-bound interpretation is relevant but the context provides less support.

The Constraint-Based framework, which guides our research, is also Context-Driven. It falls most generally into the class of generative (i.e., data explanation) approaches to perception and action (Clark, 2013) which view perception as a process of probabilistic, knowledge-driven inference. Our approach is grounded in three related observations. First, as an utterance unfolds, listeners rapidly integrate multiple sources of information. That is, utterance comprehension is probabilistic and constraint-based (MacDonald, Pearlmutter, & Seidenberg, 1994; Seidenberg & Macdonald, 1999; Tanenhaus & Trueswell, 1995). Second, listeners generate expectations of multiple types about the future, including the acoustic/phonetic properties of utterances, syntactic structures, referential domains, and possible speaker meaning (Chambers, Tanenhaus, & Magnuson, 2004; Chang, Dell, & Bock, 2006; Kutas & Hillyard, 1980; Levy, 2008; Tanenhaus, Spivey-Knowlton, Eberhard, & Sedivy, 1995; Trueswell, Tanenhaus, & Kello, 1993). Third, interlocutors can rapidly adapt their expectations to different speakers, situations, etc. (Bradlow & Bent, 2008; Clayards, Tanenhaus, Aslin, & Jacobs, 2008; Fine, Jaeger, Farmer, & Qian, 2013; Grodner & Sedivy, 2011; Kurumada, Brown, & Tanenhaus, 2012). Given these assumptions, standard hierarchical relations among different types of representations need not map onto temporal relationships in real-time processing. We illustrate these points with three examples grounded in recent results in the language processing literature.

The first example involves the mapping between speech perception, spoken word recognition, parsing, and discourse. Consider a spoken utterance that begins with “*The lamb...*” As the fricative at the onset of *the* unfolds, the acoustic signal provides probabilistic evidence for possible words, modulated by the prior expectations of these possible words, for example, their frequency, likely conditioned on the position in the utterance. Thus, after hearing only the first 50–100 ms of the signal, the listener has partial evidence that “the” is likely to be the word currently being uttered. Put differently, the hypothesis that the speaker intended to say *the* begins to provide an explanation for the observed signal. In addition, there is evidence for a likely syntactic structure (noun phrase) and evidence that the speaker intends (for some purpose) to refer to a uniquely identifiable entity (due to presuppositions associated with the definite determiner *the*, Russell, 1905). Thus, referential constraints can be available to the listener before any information in the speech signal associated with the production of the noun is available.

The second example focuses on the interpretation of referential expressions in context. Consider the instruction in (4):

(4) Put the pencil below the big apple.

Assume that the addressee is listening to the utterance in the context of a scene with a large apple, a smaller apple, a large towel, and a small pencil. According to standard accounts of scalar adjectives, a scalar dimension can only be interpreted with respect to a reference class given by the noun it modifies (e.g., what counts as big for a building differs from what counts as big for a pencil). Translating this directly into processing terms, when the listener encounters the scalar adjective *big*, interpretation should be delayed because the reference class has not yet been established. However, in practice the listener's attention will be drawn to the larger of the two apples before hearing the word *apple* because use of a scalar adjective signals a contrast among two or more entities of the same semantic type (Sedivy, Tanenhaus, Chambers, & Carlson, 1999). Thus, *apple* will be immediately interpreted as the larger of the two apples. More generally, addressees circumscribe referential domains on the fly, taking into account possible actions as well as the affordances of objects (Chambers et al., 2002; Chambers et al., 2004; Tanenhaus et al., 1995).

Finally, listeners rapidly adjust expectations about the types of utterances that speakers will produce in a particular situation. After brief exposure to a speaker who uses scalar adjectives non-canonically, for example, a speaker who frequently overmodifies, an addressee will no longer assume that a pre-nominal adjective will be used contrastively (Grodner & Sedivy, 2011). Likewise, the utterance, *It looks like a zebra* is most canonically interpreted to mean *The speaker thinks it's a zebra* (with noun focus), or *You might think it's a zebra, but it isn't* (with contrastive focus on the verb), depending on the context, the prosody, and the knowledge of the speaker and the addressee (Kurumada et al., 2012). Crucially, when a stronger alternative is sometimes used by the speaker, for example, *It is a zebra*, then the interpretation of *It looks like a zebra* with noun focus is more likely to receive the *but it's not* interpretation (Kurumada et al., 2012). Again an utterance is interpreted with respect to the context and the set of likely alternatives to the observed utterance.

When we embed the issue of when, and how quickly, upper- and lower-bound *some* are interpreted within a Constraint-Based approach with expectations and adaptation, then questions about time course no longer focus on distinguishing between the claim that scalar implicatures are computed by default and the claim that they are only computed after an initial stage of semantic processing. When there is more probabilistic support from multiple cues, listeners will compute scalar inferences more quickly and more robustly. Conversely, when there is less support, listeners will take longer to arrive at the inference, and the inference will be weaker (i.e., more easily cancelable). Under the Constraint-Based account, then, the research program becomes one of identifying the cues that listeners use in service of the broader goal of understanding the representations and processes that underlie generation of implied meanings.

1.4. Cues under investigation

We focus on two types of cues, both of which are motivated by the notion that speakers could have produced alternative utterances. The first cue is the partitive *of*, which marks the difference between (5) and (6).

(5) Alex ate some cookies.

(6) Alex ate some of the cookies.

While it may seem intuitively clear to the reader that (6) leads to a stronger implicature than (5), this intuition has not previously been tested. Moreover, researchers have used either one or the other form in their experimental stimuli without considering the effect this might have on processing. For example, some researchers who find delayed implicatures and use some form of the Literal-First hypothesis to explain these findings have consistently used the non-partitive form (Bott, Bailey, & Grodner, 2012; Bott & Noveck, 2004; Noveck & Posada, 2003). The absence of the partitive may provide weaker support for the implicature than use of the partitive form would have provided. Under a Constraint-Based account, this should result in increased processing effort.

The second cue is the availability of lexical alternatives to *some* that listeners assume are available to the speaker. In general, the alternatives to *some* will be determined by the context in which an utterance occurs. We assume that lexical items that are part of a scale will generally be available as alternatives.² However, in a given context, other lexical items may be introduced that become salient alternatives. In the current studies, we focus on the effects of number terms as alternatives, and the effects of intermixing *some* with exact number terms, for example, *You got two of the gumballs* versus *You got some of the gumballs*. The motivation for using exact number terms as a case study comes from two studies (Grodner, Klein, Carbary, & Tanenhaus, 2010; Huang & Snedeker, 2009), which used similar methods but found different results.

Both Huang and Snedeker (2009, 2011) and Grodner et al. (2010) used the visual world eye-tracking paradigm (Cooper, 1974; Tanenhaus et al., 1995). In Huang and Snedeker's (2009, 2011) experiments participants viewed a display with four quadrants, with the two left and the two right quadrants containing pictures of children of the same gender, with each child paired with objects. For example, on a sample trial, the two left quadrants might each contain a boy: one with two socks and one with nothing. The two right quadrants might each contain a girl: one with two socks (pragmatic target) and one with three soccer balls (literal target). A preamble established a context for the characters in the display. In the example, the preamble might state that a coach gave two socks to one of the boys and two socks to one of the girls, three soccer balls to the other girl, who needed the most practice, and nothing to the other boy.

Participants were asked to follow instructions such as *Point to the girl who has some of the socks*. Huang and Snedeker (2009) reasoned that if the literal interpretation is

computed prior to the inference, then, upon hearing *some*, participants should initially fixate both the semantic and pragmatic targets equally because both are consistent with the literal interpretation. If, however, the pragmatic inference is immediate, then the literal target should be rejected as soon as *some* is recognized, resulting in rapid fixation of the pragmatic target. The results strongly indicated that the literal interpretation was computed first. For commands with *all* (e.g., *Point to the girl who has all of the soccer balls*) and commands using number (e.g., *Point to the girl who has two/three of the soccer balls*), participants converged on the correct referent 200–400 ms after the quantifier. In contrast, for commands with *some*, target identification did not occur until 1,000–1,200 ms after the quantifier onset. Moreover, participants did not favor the pragmatic target prior to the noun's phonetic point of disambiguation (e.g., *-ks* of *socks*). Huang and Snedeker concluded that “even the most robust pragmatic inferences take additional time to compute” (Huang & Snedeker, 2009, p. 408).

The Huang and Snedeker results complement previous response time (Bott & Noveck, 2004; Noveck & Posada, 2003) and reading-time experiments (Breheny et al., 2006). In these studies, response times associated with the pragmatic inference are longer than both response times to a scalar item's literal meaning and to other literal controls (typically statements including *all*). Moreover, participants who interpret *some* as *some and possibly all*, so-called *logical responders* (Noveck & Posada, 2003), have faster response times than *pragmatic responders* who interpret *some* as *some but not all*.

In contrast, Grodner et al. (2010) found evidence for rapid interpretation of pragmatic *some*, using displays and a logic similar to that used by Huang and Snedeker (2009). Each trial began with three boys and three girls on opposite sides of the display and three groups of objects in the center. A pre-recorded statement described the total number and type of objects in the display. Objects were then distributed among the participants. The participant then followed a pre-recorded instruction, of the form *Click on the girl who has summa/nunna/alla/the balloons*.

Convergence on the target for utterances with *some* was just as fast as for utterances with *all*. Moreover, for trials on which participants were looking at the character with all of the objects at the onset of the quantifier *summa*) participants began to shift fixations away from that character and to the character(s) with only some of the objects (e.g., the girl with the balls or the girl with the balloons) within 200–300 ms after the onset of the quantifier. Thus, participants were immediately rejecting the literal interpretation as soon as they heard *summa*. If we compare the results for *all* and *some* in the Huang and Snedeker and Grodner et al. experiments, the time course of *all* is similar but the upper-bound interpretation of partitive *some* is computed 600–800 ms later in Huang and Snedeker.

Why might two studies as superficially similar as Huang and Snedeker (2009) and Grodner et al. (2010) find such dramatically different results? In Degen and Tanenhaus (under review), we discuss some of the primary differences between the two studies, concluding that only the presence or absence of number expressions might account for the conflicting results. Huang and Snedeker included stimuli with numbers, whereas Grodner

et al. did not. In Huang and Snedeker (2009) a display in which one of the girls had two socks was equally often paired with the instruction, *Point to the girl who has some of the socks* and *Point to the girl who has two of the socks*. In fact, Huang, Hahn, and Snedeker (2010) have shown that eliminating number instructions reduces the delay between *some* and *all* in their paradigm.

Why might instructions with exact number delay upper-bound interpretations of partitive *some*? Computation of speaker meaning takes into account what the speaker could have, but did not say, with respect to the context of the utterance. Recall that the upper-bound interpretation of *some* is licensed when the speaker could have, but did not say *all*, because *all* would have been more informative. More generally, it would be slightly odd for a speaker to use *some* when *all* would be the more natural or typical description. We propose that in situations like the Huang and Snedeker and Grodner et al. experiments, exact number is arguably more natural than *some*. In fact, intuition suggests that mixing *some* and exact number makes *some* less natural.

Consider a situation where there are two boys, two girls, four socks, three soccer balls, and four balloons. One girl is given two of the four socks, one boy the four balloons, and the other boy two of the three soccer balls. Assuming the speaker knows exactly who got what, the descriptions in (7) and (8) seem natural, compared to the description in (9):

(7) One of the girls got some of the socks and one of the boys got all of the balloons.

(8) One of the girls got two of the socks and one of the boys got all of the balloons.

(9) One of the girls got two of the socks and one of the boys got some of the soccer balls.

Grodner et al. (2010) provided some empirical support for these intuitions. They collected naturalness ratings for their displays and instructions, both with and without exact number included in the instruction set. Including exact number lowered the naturalness ratings for partitive *some* but not for *all*. However, even without exact number instructions, *some* was rated as less natural than *all*. One reason might be that in most situations, pragmatic *some* is relatively infelicitous when used to describe small sets. Again, intuition suggests that using *some* is especially odd for sets of one and two. Consider a situation where there are three soccer balls and John is given one ball and Bill two. *John got one of the soccer balls* seems a more natural description than *John got some of the soccer balls* and *Bill got two of the soccer balls* seems more natural than *Bill got some of the soccer balls*.

These observations suggest an alternative hypothesis for why responses to pragmatic *some* are delayed when intermixed with exact number for small set sizes. Most generally we suggest that *some* will compete with other rapidly available contextual alternatives. In particular, we hypothesize that the mapping between an utterance with *some* and an interpretation of that utterance is delayed when there are rapidly available, more natural alternatives to describe the state of the world. This seems likely to be the case for exact

number descriptions with small sets because the more natural number alternative is also a number in the subitizing range where number terms become rapidly available and determining the cardinality of a set does not require counting (Atkinson, Campbell, & Francis, 1976; Kaufman, Lord, Reese, & Volkman, 1949; Mandler, Shebo, & Vol, 1982). In situations where exact number is available as a description, the number term is likely to become automatically available, thus creating a more natural, more available interpretation of the scene.

In Gricean terms, we are proposing that delays in response times (to press a button or to fixate a target) which have previously been argued to be due to the costly computation of the Quantity implicature from *some* to *not all* might in fact be due to interference from lexical alternatives to *some*. In the Huang and Snedeker studies, in particular, there may be interference from number term alternatives that, while scalar in nature, function on a different scale than *some*. In Gricean terms, the motivation for this interference comes from the maxim of Manner: If there is a less ambiguous quantifier (e.g., *two*) that could have been chosen to refer to a particular set size, the speaker should have used it. If she did not, it must be because she meant something else. Finally arriving at the implicature from *some* to *not all*, when there are more natural lexical alternatives to *some* that the speaker could have used but did not, thus may involve both reasoning involving the Quantity maxim (standard scalar implicature) and the Manner maxim (inference that the speaker must have not meant the partitioned set which could have more easily and naturally been referred to by *two*). If this is indeed the case, this would have serious implications for the interpretation of response time results on scalar implicature processing across the board; previously, delays in computing pragmatic *some* have been associated with the costly computation of the scalar implicature itself, taking into account only the competition of *some* with its scalemate *all*. What we are proposing here is that this view is too narrow—rather than competing only with its lexicalized scalemate, *some* contextually competes with many other alternatives, like number terms, which are not lexicalized alternatives. Thus, observed delays in the processing of scalar implicatures might be at least partly due to costly reasoning about unnatural, misleading quantifier choices.

1.5. The gumball paradigm

The current experiments examine the role of contextual alternatives within the Constraint-Based framework using the case of exact number. Our specific hypothesis is that number selectively interferes with *some*—both in processing the upper-bound and the lower-bound interpretation—where naturalness of *some* is low and number terms are rapidly available. We evaluate this hypothesis in a series of experiments using a “gumball paradigm.” We begin with an example to illustrate how likely interpretations might change over time. Suppose that there is a gumball machine with 13 gumballs in the upper chamber. Alex knows that this gumball machine has an equal probability of dispensing 0–13 gumballs. His friend, Thomas, inserts a quarter and some number of gumballs drops to the lower chamber but Alex cannot see how many have dropped. Thomas, however, can, and he says *You got some of the gumballs*.

Before the start of the utterance Alex will have certain prior expectations about how many gumballs he got—in fact, in this case there is an equal probability of 1/14 that Alex got any number of gumballs. This is shown in the first panel of Fig. 1. Once Alex hears *You got some*, he has more information about how many gumballs he got. First, the meaning of *some* constrains the set to be larger than 0. However, Alex also has knowledge about how natural it would be for Thomas to utter *some* instead of, for example, an exact number term, to inform him of how many gumballs he got. Fig. 1 illustrates how this knowledge might shift his subjective belief probabilities for having received specific numbers of gumballs.

Alex is now more certain that he has received an intermediate set size rather than a small set (where Thomas could have easily said *one* or *two* instead of *some*) or a large set (where Thomas could have said *most*, or even *all*). Finally, once Alex hears the partitive *of*, his expectations about how many gumballs he got might shift even more (e.g., because Alex knows that the partitive is a good cue to the speaker meaning to convey upper-bound *some*), as shown in the third panel.³ Thus, by the end of the utterance Alex will be fairly certain that he did not get all of the gumballs, but he will also expect to have received an intermediate set size, as there would have been more natural alternative utterances available to pick out either end of the range of gumballs.

Note that without additional assumptions, neither the Default nor the Literal-First model makes gradient predictions about expected set size. Under the Default model, the distribution on states would be uniform until the word *some* is encountered, at which point both the zero-gumball and all-gumball state would be excluded as potential candidates, leaving a uniform distribution over states 1–12. In contrast, the Literal-First model predicts that upon encountering *some*, only the zero-gumball state should be excluded. Both models predict that over time, the integration of further contextual information may require re-including the all-state in the set of possible states (Default), or excluding the all-state from said set (Literal-First). However, neither of these models directly predicts




Fig. 1. Possible constraint-based update of most likely intended interpretation at different incremental time points. Interpretations are represented as set sizes. Bars represent the amount of probabilistic support provided for each interpretation, given the constraints at each point.

variability in naturalness of partitive versus non-partitive *some* for any set size, nor variability in the naturalness of *some* used with different set sizes. Moreover, adding probabilistic assumptions that would map onto the time course of initial processing of *some* would be inconsistent with the fundamental assumptions of each of these models.

We developed a gumball paradigm based on this scenario in order to investigate whether listeners are sensitive to the partitive and to the naturalness and availability of number descriptions as lexical alternatives to *some* in scalar implicature processing using a range of different set sizes. On each trial the participant sees a gumball machine with an upper chamber and a lower chamber, as illustrated in Fig. 2. All of the gumballs begin in the upper chamber. After a brief delay, some number of gumballs drops to the lower chamber. The participant then responds to a statement describing the scene, either by rating the statement's naturalness or judging whether they agree or disagree with the statement. We can thus obtain information about participants' judgments while at the same time recording response times as a measure of their interpretation of different quantifiers with respect to different visual scenes.

The rationale for the rating studies is that we need to establish the relative naturalness of alternative descriptions for particular set sizes. These data are crucial for generating time course predictions that distinguish the Default, Literal-First, and Constraint-Based approaches. Crucially, naturalness data are essential for evaluating the time course claims of the Constraint-Based approach. In two rating studies in which the upper chamber begins with 13 gumballs, we establish the naturalness of *some* for different set sizes of interest, in particular small sets (1–3), intermediate sets (6–8), and for the unpartitioned set (all 13 gumballs). In addition, we investigate the relative naturalness of simple *some* versus partitive *some of* for the unpartitioned set. We further investigate the effect of including exact number descriptions on naturalness ratings for *some* and *some of* used




Fig. 2. Sample displays in the gumball paradigm. Left: initial display. Right: sample second display with dropped gumballs.

with different set sizes in Experiment 2. These results allow us to make specific predictions about response times that are tested in Experiment 3.

2. Experiment 1

Experiment 1 was conducted to determine the naturalness of descriptions with *some*, *some of* (henceforth *summa*), *all of* (henceforth *all*), and *none of* (henceforth *none*) for set sizes ranging from 0 to 13.

2.1. Methods

2.1.1. Participants

Using Amazon's Mechanical Turk, 120 workers were paid \$0.30 to participate. All were native speakers of English (as per requirement) who were naïve as to the purpose of the experiment.

2.1.2. Procedure and materials

On each trial, participants saw a display of a gumball machine with an upper chamber filled with 13 gumballs and an empty lower chamber (Fig. 2). After 1.5 s, a new display was presented in which a certain number of gumballs had dropped to the lower chamber. Participants heard a pre-recorded statement of the form *You got X gumballs*, where *X* was a quantifier. They were then asked to rate how naturally the scene was described by the statement on a seven-point Likert scale, where seven was *very natural* and one was *very unnatural*. If they thought the statement was false, they were asked to click a FALSE button located beneath the scale. We varied both the size of the set in the lower chamber (0–13 gumballs) and the quantifier in the statement (*some*, *summa*, *all*, *none*). Some trials contained literally false statements. For example, participants might get none of the gumballs and hear *You got all of the gumballs*. These trials were interspersed in order to have a baseline against which to compare naturalness judgments for *some (of the)* used with the unpartitioned set. If interpreted semantically (as *You got some and possible all of the gumballs*), the *some* statement is true (however unnatural) for the unpartitioned set. However, if it is interpreted pragmatically as meaning *You got some but not all of the gumballs*, it is false and should receive a FALSE rating.

Participants were assigned to one of 24 lists. Each list contained six *some* trials, six *summa* trials, two *all* trials, and two *none* trials. To avoid an explosion of conditions, each list sampled only a subset of the full range of gumball set sizes in the lower chamber. The quantifiers *some* and *summa* occurred once each with 0 and 13 gumballs. In addition, *none* occurred once (correctly) with 0 gumballs and *all* once (correctly) with 13 gumballs. Each of *all* and *none* also occurred once with an incorrect number of gumballs. The remaining *some* and *summa* trials sampled two data points each per quantifier from the subitizing range (1–4 gumballs), one from the mid range (5–8 gumballs), and one from the high range (9–12 gumballs). For an overview, see Table 1. See Appendix A,

Table A1, for the set sizes that were sampled on each list. From each of 12 base lists, one version used a forward order and the other a reverse order. Five participants were assigned to each list.

2.2. Results and discussion

Mean ratings for each quantifier for the different set sizes are presented in Fig. 3. Clicks of the FALSE button were coded as 0. This was motivated by many participants' use of the lowest point on the scale to mean FALSE (see the histogram of responses to *some/summa* in Fig. 4). We thus treated all responses on a single, continuous dimension. Mean ratings were 5.83 for *none* (0 gumballs) and 6.28 for *all* (13 gumballs) and close to

Table 1
Distribution of the 16 experimental trials over quantifiers and set sizes

Quantifier	Set Size				
	0	Sub	Mid	High	13
Some	1	2	1	1	1
Summa	1	2	1	1	1
None	1		1		
All			1		1

Notes. See Appendix A for the exact set sizes that were sampled on different lists in the subitizing (sub), mid, and high range.




Fig. 3. Mean ratings for simple “some,” partitive “some of the,” and the exact quantifiers “none” and “all.” Means for the exact quantifiers “none” and “all” are only shown for their correct set size.




Fig. 4. Proportion of ratings (where FALSE responses are coded as 0 ratings) for “some”/“summa” in Exps. 1 and 2. Rows represent number of gumballs in the lower chamber. Note that some participants consistently gave the lowest rating on the scale (1) instead of clicking the FALSE button.

0 otherwise. Mean ratings for *some/summa* were lowest for 0 gumballs (1.4/1.1), increased to 2.7 and 5.2 for 1 and 2 gumballs, respectively, peaked in the mid range (5.81/5.73), decreased again in the high range (4.69/4.74), and decreased further at the unpartitioned set (3.47/2.7).

Our analyses were designed to address the following three questions: (a) for which set size ranges *some* and *summa* are deemed most natural; (b) whether *some* and *summa* differ in naturalness, and (c) whether the naturalness of *some* and *summa* differs for different set sizes (or ranges of set sizes).

The data were analyzed using a series of mixed effects linear regression models with by-participant random intercepts to predict ratings. *p*-values were obtained using MCMC sampling (Baayen, Davidson, & Bates, 2008). For *some/summa* used with each unique set size in the lower chamber (0–13), we fit one model each to each subset of the data corresponding to that set size in addition to trials where *all* and *none* were used with their correct set size (13 and 0, respectively). Each model included a centered fixed effect of quantifier (*some/summa* vs. *all/none*). *Some* and *summa* were most natural in the mid range (numerically peaked when used with six gumballs), where ratings did not differ from ratings for *none* and *all* used with their correct set size ($\beta = 0.2$, $SE = 0.24$, $t = 0.81$, $p < .42$). Mean ratings for *some* and *summa* did not differ for any set size except at the unpartitioned set, where *some* was more natural than *summa* ($\beta = -0.77$, $SE = 0.19$, $t = -4.07$, $p < .01$). This naturalness difference between *some* and *summa* used with the unpartitioned set suggests that *summa* is more likely to give rise to a scalar implicature than *some* and is thus dispreferred with the unpartitioned set. Because ratings for *some* and *summa* did not differ anywhere except for the unpartitioned set, we henceforth report collapsed results for *some* and *summa*.

To test the hypothesis that naturalness for *some* varies with set size, we fit a mixed effects linear model predicting mean naturalness rating from range (*subitizing* [one to four gumballs] or *mid range* [five to eight gumballs]) to the subset of the *some* cases in each range. As predicted, naturalness was lower in the subitizing range than in the mid range ($\beta = -0.79$, $SE = 0.13$, $t = -6.01$, $p < .001$). Similarly, naturalness ratings for *some* were lower for the unpartitioned set than in the mid range ($\beta = -2.68$, $SE = 0.15$, $t = -17.82$, $p < .001$). However, Fig. 3 suggests that the naturalness ratings differed for set sizes within the subitizing range (1–4). Performing the analysis on subsets of the data comparing each set size with naturalness of *some/summa* used with the preferred set size (six gumballs) yields the following results. Collapsing over *some* and *summa*, we coded small set versus six gumballs as 0 and 1, respectively, and subsequently centered the predictor. The strongest effect is observed for one gumball ($\beta = 2.95$, $SE = 0.17$, $t = 14.28$, $p < .0001$). The effect is somewhat weaker for two ($\beta = 1.0$, $SE = 0.22$, $t = 4.47$, $p < .0001$), even weaker for three ($\beta = 0.36$, $SE = 0.2$, $t = 1.81$, $p = .05$), and only marginally significant for four ($\beta = 0.29$, $SE = 0.19$, $t = 1.52$, $p < .1$).⁴ The coefficients for the set size predictor for each subset model are plotted in Fig. 5. Given these results we will refer to “small set size” effects rather than “subitizing” effects.

In sum, *some* and *summa* were both judged to be quite unnatural for set sizes of 1 and 2, and more natural but not quite as natural as for the preferred set size (six gumballs) for 3. Naturalness also decreased after the mid range (five to eight gumballs) and was low at the unpartitioned set. In addition, the partitive, *some of*, was less natural to refer to the unpartitioned set than simple *some*.

Finally, we note that naturalness ratings for *some/summa* gradually decreased for set sizes above 6. This is probably due to there being other more natural,




Fig. 5. Set size model coefficient for each set size in the subitizing range. Error bars represent 1 SE.

salient alternatives for that range: *many* and *most*. It is striking that these alternatives seem to affect the naturalness of *some* just as much as number terms for small set sizes.

The naturalness results from this study point to an interesting fact about the meaning of *some*. The linguistic literature standardly treats the semantics of *some* as proposed in Generalized Quantifier Theory (Barwise & Cooper, 1981) as the existential operator (corresponding to *at least one*). Under this view, as long as at least one gumball dropped to the lower chamber, participants should have rated the *some* statements as true (i.e., not clicked the FALSE button). However, this was not the case: *some* received 12% FALSE ratings for one gumball and 9% FALSE ratings for the unpartitioned set; *summa* statements were rated FALSE in 7% of cases for one gumball and 18% of cases for the unpartitioned set. For comparison, rates of FALSE responses to *some/summa* for all other correct set sizes were 0%.

In addition, treating *some* as simply the existential operator does not allow a role for the naturalness of quantifiers. What matters is that a statement with *some* is true or false. Differences in naturalness are not predicted. Whether this means that language users' underlying representation of *some* is more complex than the existential operator (and similarly for other quantifiers) is an open question. One could argue for an analogy to the distinction between the underlying category and what affects categorization (for discussion of this perspective, see Armstrong, Gleitman, & Gleitman, 1983). However, our preferred view is that for the purposes of formalizing truth conditions, the existential operator is useful as an abstraction over the possible contexts in which a simple statement with *some* could be true. The underlying cognitive representations, on the other hand, are likely to involve mappings onto expectations of usage in specific contexts.

One way of conceptualizing these naturalness results is that we have obtained probability distributions over set sizes for different quantifiers, where the relevant probabilities are participants' subjective probabilities of expecting a speaker to use a particular quantifier with a particular set size. Thus, a vague quantifier like *some*, where naturalness is high for intermediate sets and gradually drops off at both ends of the spectrum, has a very wide distribution, with probability mass distributed over many set sizes. In contrast, for a number term like *two*, one would expect naturalness to be very high for a set size of two and close to 0 for all other cardinalities, and thus the distribution would be very narrow and peaked around 2.

According to the Constraint-Based account that allows for parallel processing of multiple sources of information, distributions of quantifiers over set sizes (or in other words listeners' expectations about speakers' quantifier use) are a function of at least two factors: (a) set size and (b) awareness of contextual availability of alternative quantifiers. If no lexical alternative is available, listeners will have some expectations about the use of *some* with different set sizes. We propose that the distribution of ratings obtained in Experiment 1 reflects just these expectations. For *some*, naturalness is highest for intermediate set sizes and drops off at both ends of the tested range. That is, listeners' expectation for *some* to be used is highest in the mid range. The Constraint-Based account predicts that listeners' expectations about quantifier use are sensitive to alternatives. Including number terms among the experimental items, thus making participants aware that number terms are contextually available alternatives to *some*, should change this distribution. In particular, the prediction is that due to subitizing processes, which allow number terms to become rapidly available as labels for small sets, the naturalness of *some* should decrease for small sets when number terms are included. In other words, participants' expectations that a small set will be referred to by *some* should decrease. This prediction is tested in Experiment 2.

3. Experiment 2

Experiment 2 tested the hypothesis that when number terms are included as alternatives within the context of an experiment, the naturalness of *some* will be reduced when it is used with small set sizes. Using the same paradigm as in Experiment 1, we included number terms among the stimuli to test the hypothesis that the naturalness of *some/summa* would be reduced when used with small set sizes, where number terms are hypothesized to be most natural.

3.1. Methods

3.1.1. Participants

Using Amazon's Mechanical Turk, 240 workers were paid \$0.75 to participate. All were native speakers of English (as per requirement) who were naïve as to the purpose of the experiment.

3.1.2. Procedure and materials

The procedure was the same as that described for Experiment 1 with one difference; the number terms *one of the* through *twelve of the* were included among the stimuli. Each participant rated naturalness of statements with quantifiers as descriptions of gumball machine scenes on 32 trials. Participants were assigned to one of 48 lists. As in Experiment 1, each list contained six *some* trials, six *summa* trials, two *all* trials, and two *none* trials (see Table 1 for an overview of these 16 trials). In addition, four number terms were included on each list. Each number term occurred once with its correct set size, once with a wrong set size that differed from the correct set size by one gumball, and once with a wrong set size that differed from the correct set size by at least three gumballs. The lists were created from the same base lists used in Experiment 1. See Appendix A for the set sizes that were sampled on each list. Four versions of each of the twelve base lists were created. On half of the lists, *some/summa* occurred before the correct number term for each set size; on the other half, it occurred after the correct number term. Half of the lists sampled wrong set sizes that were one bigger than the correct set size for the number terms employed, and half sampled set sizes that were one smaller.

3.2. Results

As in Experiment 1, clicks of the FALSE button were coded as 0 and ratings treated as continuous values (see Fig. 4 for a histogram of the distribution of ratings for each number of gumballs). Mean ratings were 5.71 for *none* with 0 gumballs and 6.31 for *all* with 13 gumballs, and close to 0 otherwise. Mean ratings for *some/summa* were lowest for 0 gumballs (1.08/0.96), increased from 2.02/1.99 to 4.91/4.54 in the small set range, peaked in the mid range at six gumballs (5.82/5.97), decreased again in the high range (4.33/4.47), and decreased further at the unpartitioned set (3.42/2.65), replicating the general shape of the curve obtained in Experiment 1.

Again, *some* and *summa* were most natural when used with six gumballs. As in Experiment 1, mean ratings for *some* and *summa* did not differ for any set size except at the unpartitioned set, where *some* was more natural than *summa* ($\beta = -0.77$, $SE = 0.13$, $t = -4.07$, $p < .001$).

To test the hypothesis that adding number terms decreases the naturalness for *some/summa* in the small set range but nowhere else, we fit a series of mixed effects linear models with random by-participant intercepts, predicting mean naturalness rating from number term presence for each range (no gumballs, small, mid, high, unpartitioned set). The models were fit to the combined data sets from Experiment 1 (numbers absent) and Experiment 2 (numbers present). As predicted, naturalness was lower for both *some* and *summa* when numbers were present in the small set range ($\beta = -0.49$, $SE = 0.16$, $t = -3.13$, $p < .002$), but not for 0 gumballs ($\beta = -0.19$, $SE = 0.17$, $t = -1.11$, $p < .28$), in the mid range ($\beta = -0.14$, $SE = 0.14$, -0.99 , $p = .13$), in the high range ($\beta = -0.38$, $SE = 0.23$, $t = -1.37$, $p < .12$), or with the unpartitioned set ($\beta = -0.11$, $SE = 0.23$, $t = -0.49$, $p < .71$). Fig. 6 presents the mean naturalness ratings when numbers were present (Experiment 2) and absent (Experiment 1).




Fig. 6. Mean ratings for “some” (collapsing over simple and partitive “some”) and exact quantifiers/number terms when number terms are present (Experiment 2) versus absent (Experiment 1). Means for the exact quantifiers are only shown for their correct set size.

Performing each analysis individually for each set size in the subitizing range shows that the strength of the number presence effect in the small set range differed for different set sizes. It was strongest for one gumball ($\beta = -0.83$, $SE = 0.27$, $t = -3.12$, $p < .0001$), less strong for two ($\beta = -0.54$, $SE = 0.22$, $t = -2.4$, $p < .01$), trending for three gumballs ($\beta = -0.27$, $SE = 0.2$, $t = -1.16$, $p < .12$), and nonsignificant for four gumballs ($\beta = -0.1$, $SE = 0.18$, $t = -0.57$, $p < .61$). We provide a coefficient plot for the effect of number presence for different set sizes in Fig. 7. Therefore, naturalness effects are not due to subitizing *per se*, as we had initially hypothesized. Rather, subitizing might interact with naturalness to determine the degree to which number alternatives compete with *some*.

Number terms did not reduce naturalness for *none* ($\beta = -0.1$, $SE = 0.22$, $t = -0.47$, $p < .6$) and *all* ($\beta = -0.02$, $SE = 0.17$, $t = 0.14$, $p < .84$) when used with their correct set sizes. Finally, although ratings for number term ratings are extremely high throughout when used with their correct set size and close to floor otherwise, number terms are judged as more natural when used with small set sizes than when used with large ones as determined in a model predicting mean naturalness rating from a continuous set size predictor ($\beta = -0.05$, $SE = 0.01$, $p < .001$). The exception is *ten*, which is judged to be slightly more natural than the surrounding terms.

3.3. Discussion

The results of the two rating studies suggest that a listener’s perception of an expression’s naturalness is directly affected by the availability of lexical alternatives. With the




Fig. 7. Model coefficients for number term presence predictor for each set size in the subitizing range. Error bars represent 1 SE.

exception of 6 and 7 (around half of the original set size), numbers are always judged to be more natural than *some/summa* when they are intermixed. This difference, however, is largest for the smallest set sizes. As predicted, the reduced naturalness of *some/summa* used with small sets, established in Experiment 1, decreased further when number terms were added to the stimuli. Therefore, at least in off-line judgments, listeners take into account what the speaker could have said, but did not. The results of Experiments 1 and 2 establish that the naturalness of descriptions with *some* varies with set size and for small set sizes is affected by the inclusion of number. Note, as mentioned in the Introduction, that these patterns are not predicted by the Literal-First and Default models. However, because these models focus on differences in the time-course of processing, these results cannot be taken as evidence against these models. Crucially, we use the obtained naturalness ratings to test the Default, Literal-First, and Constraint-Based model in Experiment 3, which evaluated competing predictions about time course using response times as the primary dependent measure.

4. Experiment 3

Experiment 3 was designed to test whether the effect of available natural alternatives is reflected in response times. Using the same paradigm and stimuli, we recorded participants' judgments and response times to press one of two buttons (YES or NO) depending on whether they agreed or disagreed with the description. Based on the naturalness results from Experiments 1 and 2, the Constraint-Based account predicts that participants' YES

responses should be slower for more unnatural statements. Specifically, for *some* and *summa* response times are predicted to be slower compared to their more natural alternatives when used with (a) the unpartitioned set, where *all* is a more natural alternative and (b) in the small set range, where number terms are more natural and more rapidly available. Based on the naturalness data, the largest effect is expected for a set size of one, a somewhat smaller effect for two, and a still smaller effect for three. Response times for *some/summa* with these set sizes should be slower than when *some* and *summa* are used in the preferred range (four to seven gumballs).

In Experiments 1 and 2, we also observed a difference in naturalness of simple versus partitive *some* for the unpartitioned set. This difference should be reflected both in the number of YES responses (more YES responses to *some* than to *summa*) and in response times (faster YES responses for *some* than for *summa*).

Note that neither the Default nor the Literal-First model predicts response time differences based on naturalness of alternatives—regardless of set size, processing of a statement with *some* should take the same amount of time, except for the unpartitioned set, where the Default model predicts longer response times for semantic YES responses and the Literal-First model predicts longer response times for pragmatic NO responses. In addition, neither of these models predicts when a statement with *some* should result in a pragmatic NO judgment despite being semantically true. In contrast, the Constraint-Based account predicts the proportion of NO judgments to be proportional to the naturalness of *some* used with that set size.

The conditions in which *some/summa* are used with the unpartitioned set are of additional interest because they can be linked to the literature using sentence verification tasks. In these conditions, enriching the statement to *You got some but not all (of the) gumballs* via scalar implicature makes it false. However, if no such pragmatic enrichment takes place, it is true. That is, YES responses reflect the semantic, *at least*, interpretation of the quantifier, whereas NO responses reflect the pragmatic, *but not all*, interpretation. Noveck and Posada (2003) and Bott and Noveck (2004) called the former *logical responses* and the latter *pragmatic responses*; we will make the same distinction but use the terms *semantic* and *pragmatic* and in addition characterize responders by how consistent they were in their responses. Analysis of semantic and pragmatic response times will allow us to test the predictions of the Constraint-Based, Literal-First, and Default account.

In Bott and Noveck's sentence verification paradigm, participants were asked to perform a two-alternative, forced-choice task. Participants were asked to respond TRUE or FALSE to clearly true, clearly false, and underinformative items (e.g., *Some elephants have trunks*). Bott and Noveck found that (a) pragmatic responses reflecting the implicature were slower than semantic responses; and (b) pragmatic responses were slower than TRUE responses to *all* for the unpartitioned set. If processing of the scalar item *some* proceeds similarly in our paradigm, we would expect to replicate Bott and Noveck's pattern of results for the unpartitioned set with YES responses to *some* being faster than NO responses, and NO responses to *some* being slower than YES responses to *all*.

[illegible]

gumballs), and one of 0 or 13 gumballs. Finally, four additional number terms were sampled (one each) from the set of *three* or *four*, *five* to *seven*, *eight* or *nine*, and *ten* to *twelve*. This ensured that number terms were not all clustered at one end of the full range. Each number term occurred four times with its correct set size and four times with an incorrect number, one each sampled from the small set range, the mid range, the high range, and one of 0 or 13 gumballs, excluding the correct set size. For example, *three* could occur 4 times with 3 gumballs, once with 4, once with 7, once with 11, and once with 13 gumballs. The reason we included so many false number trials was to provide an approximate balance of YES and NO responses to avoid inducing an overall YES bias that might influence participants' response times.

To summarize, there were 28 *some* and *summa* trials each, 16 *all* trials, 16 *none* trials, 8 *one* trials, 8 *two* trials, and 8 trials each for four additional number terms. Of these, 64 were YES trials, 60 were NO trials, and 12 were critical trials—cases of *some/summa* used with the unpartitioned set, where they were underinformative, and *some/summa* used with one gumball, where a NO response is expected if the statement triggers a plural implicature (*at least two gumballs*, Zweig, 2009) and a YES response if it does not. Finally, there were three different versions of each image, with slightly different arrangements of gumballs to discourage participants from forming quantifier—image associations.

4.2. Results

A total of 6,392 responses were recorded. Of those, 26 trials were excluded because participants did not respond within the 4 seconds provided before the trial timed out. These were mostly cases of high number terms occurring with big set sizes that required counting (e.g., 11 *eleven* trials, 8 *twelve* trials, 5 *ten* trials). Further, 33 cases with response times above or below 3 *SDs* from the grand mean of response times were also excluded. Finally, 254 cases of incorrect responses were excluded from the analysis. These were mostly cases of quantifier and set size combinations where counting a large set was necessary and the set size differed only slightly from the correct set size (e.g., *ten* used with a set size of 9). In total, 4.9% of the initial data set was excluded from the analysis.

We organize the results as follows. We first report the proportion of YES and NO responses, focusing on the relationship between response choice and the naturalness ratings. These judgments will not allow us to tease apart the predictions of the Default and Literal-First model because these models do not make any predictions about response choices. However, they will allow us to test whether the naturalness differences between *all*, *some*, and *summa* are reflected in participants' binary response choices.

We then turn to the relationship between response times and naturalness ratings, testing the predictions we outlined earlier. Finally, we examine judgments and response times for pragmatic and semantic responses, relating our results to earlier work by Bott and Noveck (2004) and Noveck and Posada (2003). The results will be discussed with respect to the predictions made by the Default, Literal-First, and Constraint-Based models.

4.2.1. Judgments

All statistical analyses were obtained from mixed effects logistic regressions predicting the binary response outcome (YES or NO) from the predictors of interest. All models were fitted with the maximal random effects structure with by-participant random slopes for all within-participant factors of interest unless mentioned otherwise, following the guidelines in Barr, Levy, Scheepers, and Tily (2013).

We first examine the unpartitioned *some/summa* conditions, which are functionally equivalent to the underinformative conditions from Noveck and Posada (2003) and Bott and Noveck (2004). Recall that under a semantic interpretation of *some* (*You got some, and possibly all of the gumballs*), participants should respond YES when they get all of the gumballs, while a pragmatic interpretation (*You got some, but not all of the gumballs*) yields a NO response. The judgment data qualitatively replicate the findings from the earlier studies: 100% of participants' responses to *all* were YES, compared with 71% YES responses to partitive *some*. Judgments for simple *some* were intermediate between the two, with 82% YES responses. The difference between *some* and *summa* was significant in a mixed effects logistic regression predicting the binary response outcome (YES or NO) from quantifier (*some* or *summa*). The log odds of a YES response are lower for *summa* than *some* ($\beta = -1.18$, $SE = 0.34$, $p < .001$), reflecting the naturalness results obtained in Experiments 1 and 2, where *summa* was judged as less natural than *some* when used with the unpartitioned set. Thus, both the word *some* and its use in the partitive construction increase the probability and/or strength of generating an implicature, consistent with the naturalness-based predictions of the Constraint-Based account.

4.2.2. Response time analysis of naturalness effects on YES responses

Response times ranged from 577 to 3,574 ms (mean: 1,420 ms, SD : 444 ms). Results of statistical analyses were obtained from mixed effects linear regression models predicting log-transformed response times from the predictors of interest. As with the judgment data, all models were fitted with the maximal random effects structure with by-participant random slopes for all within-participant factors of interest. Significance of predictors was confirmed by performing likelihood ratio tests, in which the deviance ($-2LL$) of a model containing the fixed effect is compared to another model without that effect that is otherwise identical. This is one of the procedures recommended by Barr et al. (2013) for models containing random correlation parameters, as MCMC sampling (the approach recommended by Baayen et al., 2008) is not implemented in the R lme4 package.

For YES responses at the unpartitioned set, quantifier was Helmert-coded. Two Helmert contrasts over the three levels of quantifier were included in the model, comparing each of the more natural levels against all less natural levels (*all* vs. *some/summa*, *some* vs. *summa*). YES responses to *all* were faster than *some/summa* ($\beta = -0.26$, $SE = 0.02$, $t = -10.79$) and YES responses to *some* were faster than to *summa* ($\beta = -0.09$, $SE = 0.03$, $t = -3.24$). YES responses to *some* and *summa* were slower for the unpartitioned set than in the most natural range determined in Experiments 1 and 2, four to seven gumballs ($\beta = 0.09$, $SE = 0.03$, $t = 3.44$).

A similar pattern holds in the small set range for the comparison between *some/summa* and number terms: Responses to both *some* ($\beta = 0.12$, $SE = 0.02$, $t = 5.7$) and *summa* ($\beta = 0.12$, $SE = 0.02$, $t = 5.8$) were slower than number terms. Response times in the small set range did not differ for *some* and *summa*, as determined in model comparison between a model with and without a quantifier predictor ($\beta = 0.02$, $SE = 0.02$, $t = 1.4$), so we collapse them in further analysis.

There was a main effect of set size in the small set range: Responses were faster as set size increased ($\beta = -0.02$, $SE = 0.009$, $t = -2.54$). The interaction between set size and quantifier (number term vs. *some*) was also significant ($\beta = -0.08$, $SE = 0.02$, $t = -4.32$), such that there was no difference in response times for number terms used with different set sizes in the subitizing range, but response times decreased for *some/summa* with increasing set size. That is, the difference in response time between *some/summa* and number terms is largest for one gumball, somewhat smaller for two gumballs, and smaller still for three gumballs. This mirrors the naturalness data obtained in Experiments 1 and 2 (see Fig. 6). Comparing response times for *some/summa* in the small set range to those in the preferred range (4–7), the results are similar: Responses in the small set range are slower than in the preferred range ($\beta = 0.04$, $SE = 0.02$, $t = 2.3$). Mean response times for YES responses are shown in Fig. 8 (response times for *some* and *summa* are collapsed as they did not differ).

Analyzing the overall effect of naturalness on response times (not restricted to *some/summa*) yields the following results. The Spearman r between log-transformed response times and mean naturalness for a given quantifier and set size combination was -0.1 for




Fig. 8. Mean response times of YES responses to “some” (collapsed over simple and partitive use) and for exact quantifiers and number terms. For exact quantifiers, only the response time for their correct cardinality is plotted.

YES responses overall (collapsed over quantifier). This value increased to -0.3 upon exclusion of cases of number terms used outside the subitizing range, where counting is necessary to determine set size. This correlation was significant in a model predicting log-transformed response time from a centered naturalness predictor and a centered control predictor coding cases of number terms used outside the small set range as 1 and all other cases as 0. The main effect of naturalness was significant in the predicted direction, such that more natural combinations of quantifiers and set sizes were responded to more quickly than less natural ones ($\beta = -0.04$, $SE = 0.01$, $t = -9.3$). In addition, a main effect of the control predictor revealed that number terms used outside the subitizing range are responded to more slowly than cases that do not require counting ($\beta = 0.44$, $SE = 0.02$, $t = 18.6$).

4.2.3. Judgments and YES response times analyzed by pragmatic and semantic responders

Table 3 shows the distribution of participants over number of semantic responses to *some/summa* at the unpartitioned set. Noveck and Posada (2003) and Bott and Noveck (2004) found that individual participants had a strong tendency to give mostly pragmatic or mostly semantic responses at the unpartitioned set.⁵ Therefore, they conducted subanalyses comparing pragmatic and semantic responders. Our participants were less consistent. Rather than two groups of responders clustered at either end of the full range, we observe a continuum in participants' response behavior, with more participants clustered at the semantic end. Forty-two percent of the participants gave 100% (8) semantic responses. Dividing participants into two groups, semantic and pragmatic responders, where pragmatic responders are defined as those participants who gave pragmatic responses more than half of the time and semantic responders as those who responded semantically more than half of the time yields a large group of semantic responders (38 participants, 81%) and a smaller group of pragmatic responders (7 participants, 15%). Two participants (4%) gave an equal number of semantic and pragmatic responses.

Given the nature of the distribution, rather than analyzing response times for semantic and pragmatic responders separately, we included a continuous predictor of responder type (degree of "semanticity" as determined by number of semantic responses) as a control variable in the analyses.⁶ That is, a participant with one semantic response was treated as a more pragmatic responder than a participant with five semantic responses. We analyzed the effect of (continuous) responder type on the response time effects reported above, specifically (a) the naturalness effect at the unpartitioned set and (b) the naturalness effect for small sets. First, we included centered continuous responder type as interaction terms with the Helmert contrasts for quantifier (*all* vs. *some/summa*, *some* vs. *summa*) for the unpartitioned set. In this model, both interactions were significant: the

Table 3
Distribution of participants over number of semantic responses given

Number of semantic responses	0	1	2	3	4	5	6	7	8
Number of participants	2	2	2	1	2	5	6	6	21

difference between *some* and *summa* was more pronounced for more pragmatic responders ($\beta = 0.05$, $SE = 0.02$, $t = 2.5$), while the difference between *all* and *some/summa* was significantly different for different responder types ($\beta = 0.04$, $SE = 0.01$, $t = 3.97$) but seems to be better accounted for by participants' response consistency (see below). This suggests that more pragmatic responders are more sensitive to the relative naturalness of simple and partitive *some* used with an unpartitioned set. We return to this finding in the discussion.

An analysis of responder type for the naturalness effects for small set sizes yielded no significant results. That is, responder type did not interact with the quantifier by set size interaction reported above.

4.2.4. Response time analysis for semantic versus pragmatic responses to *some/summa*

Recall that the Default model predicts pragmatic NO responses to *some/summa* with the unpartitioned set to be faster than semantic YES responses, while the reverse is the case for the Literal-First hypothesis. The latter has found support in a similar sentence verification task as the one reported here (Bott & Noveck, 2004). We thus attempted to replicate Bott and Noveck's finding that pragmatic NO responses are slower than semantic YES responses. To this end we conducted four different analyses. All were linear regression models predicting log-transformed response time from response and quantifier predictors and their interaction (the interaction terms were included to test for whether NO and YES responses were arrived at with different speed for *some* and *summa*): In model 1, we compared all YES responses to all NO responses. Then we performed the Bott and Noveck between-participants analysis comparing only responses from participants who responded entirely consistently to *some/summa* (i.e., either 8 or 0 semantic responses in total) in model 2. In a very similar between-participants analysis, we compared response times from participants who responded entirely consistently to either *some* or *summa* (i.e., either 4 or 0 semantic responses to either quantifier) in model 3. Finally, again following Bott and Noveck, we compared response times to YES and NO responses within participants, excluding the consistent responders that entered model 2 from model 4. Results are summarized in Table 4.

The interaction between quantifier and response was not significant in any of the models, suggesting there was no difference between *some* and *summa* in the speed with which participants responded YES or NO. The main effect of quantifier reached significance in both model 1 (all responses to *some/summa* at unpartitioned set) and model 4 (including only inconsistent responders), such that responses to *summa* were generally slower than those to *some* (see Table 4 for coefficients). Finally, the main effect of response was marginally significant in models 2 and 3 (including only consistent responders, either overall, or within quantifier condition), such that YES responses were marginally faster than NO responses.

4.2.5. Response times as a function of response inconsistency

We conducted a final response time analysis that was motivated by the overall inconsistency in participants' response behavior at the unpartitioned set. Rather than analyzing

Table 4

Model coefficients, standard error, *t*-value, and *p*-value for the three predictors (quantifier, response, and their interaction) in each of four different models

Model	Obs.	Quantifier (<i>some</i> , <i>summa</i>)				Response (<i>no</i> , <i>yes</i>)				Quantifier : Response Interaction			
		β	<i>SE</i>	<i>t</i>	<i>p</i>	β	<i>SE</i>	<i>t</i>	<i>p</i>	β	<i>SE</i>	<i>t</i>	<i>p</i>
1. Overall	375	0.06	0.02	2.69	<.05	0.01	0.04	0.25	.86	−0.05	0.06	−0.88	.23
2. Consistent	184	0.03	0.03	0.95	.38	−0.19	0.15	−1.25	<.09	0.23	0.1	0.24	.83
3. Consistent within quantifier	256	0.01	0.03	0.50	.92	−0.12	0.08	−1.51	<.06	−0.02	0.09	−0.23	.52
4. Inconsistent	191	0.1	0.04	2.78	<.05	0.03	0.04	0.76	.53	−0.04	0.08	−0.56	.33

Note. Overall: model comparing all YES responses to all NO responses. Consistent: model comparing YES and NO responses of completely consistent responders. Consistent within quantifier: model comparing YES and NO responses of participants who gave completely consistent responses within either the *some* or the *summa* condition at the unpartitioned set. Inconsistent: model comparing YES and NO responses of participants who gave at least one inconsistent response.

only how participants' degree of semanticity impacted their response times, we analyzed the effect of within-participant response inconsistency on response times. Five levels of inconsistency were derived from the number of semantic responses given. Participants with completely inconsistent responses (four semantic and four pragmatic responses) were assigned the highest inconsistency level (5). Participants with a 3:5 or 5:3 distribution were assigned level 4, a 2:6, or 6:2 distribution were assigned level 3, a 1:7/7:1 distribution level 2, and a 0:8/8:0 distribution (participants who gave only semantic or only pragmatic responses) level 1.

There is a clear non-linear effect of inconsistency on YES responses for the unpartitioned set (see Fig. 9).

Model comparison of models with polynomial terms of different orders for inconsistency and their interactions with naturalness reveal that there is a significant main effect of naturalness (as observed before, $\beta = -0.07$, $SE = 0.01$, $t = -11.0$), a significant effect of the second-order inconsistency term ($\beta = 1.62$, $SE = 0.6$, $t = 2.7$), and a significant interaction of naturalness and the first-order inconsistency term ($\beta = -0.65$, $SE = 0.17$, $t = -3.78$). That is, those participants who responded most inconsistently also responded most slowly to *some*, *summa*, and *all* when used with the unpartitioned set. Response times decreased as inconsistency increased, but both completely pragmatic and completely semantic responders showed a slight but significant increase in response times (as revealed in the significant second-order term).

As we did for responder type, we analyzed the effect of inconsistency on the naturalness effect at the unpartitioned set by including centered inconsistency level as interaction terms with the Helmert contrasts for quantifier (*all* vs. *some/summa*, *some* vs. *summa*).

This yields a strong interaction of inconsistency with the *all* versus *some/summa* contrast ($\beta = -0.08$, $SE = 0.02$, $t = -3.36$) and a weak trend for the interaction with the *some* versus *summa* contrast ($\beta = -0.06$, $SE = 0.03$, $t = 1.42$) in the same direction. The

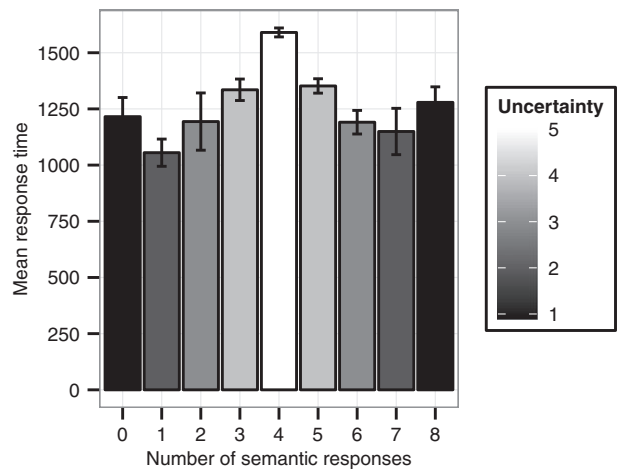


Fig. 9. Mean response times to YES responses at the unpartitioned set (collapsed across quantifiers) as a function of the degree of each responder’s response semanticity.




Fig. 10. Model coefficients obtained from models predicting YES response times from Helmert-coded quantifier contrasts at the unpartitioned set. Coefficients are plotted as a function of responder inconsistency. Numbers below the bars indicate that coefficient’s *p*-value.

effect of the difference between the different levels of the quantifier predictor on log response times is larger for more inconsistent responders. This is visualized in Fig. 10, where the coefficients of the two contrasts are plotted for models run on subsets of the data depending on inconsistency level. As with responder type, letting the inconsistency term interact with quantifier (*some/summa* or *number*) for small sets replicates the effects reported previously, but there are no significant interactions with inconsistency.

4.3. Discussion

In Experiment 3, we found that the mean naturalness of the quantifier *some* for different set sizes predicted response times. In particular, YES responses were processed more slowly with decreasing naturalness. This effect was particularly strong where the more natural alternative was very rapidly available due to subitizing processes (e.g., *two*). In addition, responders who gave many pragmatic responses to *some/summa* were more sensitive to the naturalness difference between *some* and *summa* when used with the unpartitioned set, whereas less consistent participants were more sensitive to the naturalness difference between *all* and *some/summa*. The main result of Bott and Noveck's, namely that pragmatic NO responses to *some/summa* are slower than semantic ones, was replicated marginally in a subset of analyses we conducted. We discuss both the responder type/inconsistency effects and the implications of the partial replication of Bott and Noveck's results in more detail here and defer discussion of the more general theoretical implications until the General discussion.

4.3.1. Response inconsistency as uncertainty about the QUD

The difference in naturalness effects for pragmatic and semantic responders raises two questions. First, why do more pragmatic responders exhibit a stronger naturalness effect than more semantic responders? Second, why do more inconsistent responders exhibit a stronger sensitivity to the difference between *all* and *some/summa*? Here we draw upon the notion of the QUD (Roberts, 1996), which is becoming increasingly influential within formal pragmatics. The QUD is the question that a participant in a discourse is trying to answer at the point of an utterance. Depending on what the QUD is, an utterance of the same sentence in different contexts may lead to different implicatures. Zondervan (2008) has shown that the QUD can affect how reliably an implicature is drawn: More exclusivity implicatures were triggered by the scalar item *or* when it was in a focused constituent.

Many researchers have noted that scalar implicatures arise only when the stronger alternative is contextually relevant (e.g., Carston, 1998; Green, 1995; Levinson, 2000). Thus, for example, (6) in response to (5a) is taken to implicate that not all of their documents are forgeries, whereas the same implicature does not arise if (6) is uttered in response to (5b) (Levinson, 2000).

(5a) Are all of their documents forgeries?

(5b) Is there any evidence against them?

(6) Some of their documents are forgeries.

Similarly, there are many different potential QUDs that the utterance *You got some of the gumballs* might be interpreted relative to, as is illustrated in the examples in (7):

(7a) Did I get all of the gumballs?

(7b) How many gumballs did I get?

(7c) Did I get any of the gumballs?

Intuitively, *You got some of the gumballs* implicates that you did not get all of them if it is an answer to the question in (7a), but not to the question in (7c), and more weakly to the question in (7b). This is compatible with the probabilistic account of scalar implicatures proposed by Benjamin Russell (2012), where the strength of a scalar implicature depends on the relative relevance to the QUD of the stronger and weaker alternatives. For space reasons we cannot work this out formally here, but the relative difference in relevance between the *all* and the *some* alternative is greatest for the QUD in (7a) (where the *all* alternative completely resolves the QUD but the *some* alternative does not), smallest for the QUD in (7c) (where both the *some* and the *all* alternative resolve the QUD), and intermediate for (7b) (where *all* resolves the QUD and *some* at least reduces the uncertainty about the number of gumballs received).

In a particular context, the QUD is sometimes established explicitly, for example, by asking a question. However, the QUD is often a matter of negotiation between interlocutors. That is, at a particular point in discourse interlocutors may have *uncertainty* about the actual QUD. One way to view the inconsistency results above is that participants in our experiment had different amounts of uncertainty about the QUD when they observed an utterance of *You got some (of the) gumballs* with the unpartitioned set. Under this view, participants who consistently responded either semantically or pragmatically had little uncertainty about the QUD: Semantic responders consistently adopted the QUD in (7c), whereas pragmatic responders consistently adopted the QUD in (7a). The distributions of participants' responses who were intermediate between the two extremes reflect their increased uncertainty about the actual QUD and it may have been this uncertainty that resulted in slower response times. This explanation is also consistent with the result that more inconsistent/uncertain responders are more sensitive to the difference between *all* and *some/summa* than to the difference between *some* and *summa*: *You got all of the gumballs* is true for the unpartitioned set regardless of the QUD; thus, the difference between more and less certain responders should be small in their response times to *all*, and it was. Uncertainty about the QUD should have a much larger effect on responses to *some/summa*, as the truth or falsity of *You got some (of the) gumballs* with the unpartitioned set depends on the QUD. Greater uncertainty about the QUD should result in slower responses. This is just the pattern that we observed.

We acknowledge that our appeal to the uncertainty of the QUD to explain response time patterns is necessarily post-hoc because we did not manipulate the QUD. In future research, it will be important to investigate the effect of QUD uncertainty on response times by explicitly manipulating the (implicit or explicit) QUD that statements containing scalar items are interpreted relative to. We are currently developing paradigms to pursue this line of inquiry.

The second interesting difference in responder type was that participants who gave more pragmatic responses to *some/summa* at the unpartitioned set were more sensitive to

the naturalness difference between *some* and *summa*; when compared to semantic responders, they responded YES more slowly to *summa* relative to *some*.

It is possible that there was a difference in our participant population between listeners who are generally more sensitive to naturalness differences between different utterance alternatives and those less sensitive to naturalness differences. The more sensitive ones should have responded more pragmatically overall because the naturalness of *some/summa* is generally rated lower than the naturalness of *all* for the unpartitioned set. Therefore, participants who are more sensitive to naturalness are likely to make more pragmatic responses. It is plausible to assume that participants who are more sensitive to the naturalness difference between *all* and *some/summa* with the unpartitioned set (as reflected in judgments) are also more sensitive to the difference between *some* and *summa*. If this difference as reflected in response times, it would result in the pattern of results we observe.

4.3.2. Implications of the response time analysis of pragmatic versus semantic judgments

Recall that Bott and Noveck (2004) found that pragmatic responses to *some* were slower than semantic responses both between and within participants. They interpreted this as evidence that the interpretation of the semantic *some* and possibly *all* interpretation comes for free and is more basic than the *some but not all* interpretation, which is assumed to be computed only if the semantic interpretation does not suffice to meet expectations of relevance in context and incurs additional processing cost.

We partly replicated this result: In the analysis comparing semantic and pragmatic responses between participants who were entirely consistent in their judgments to *some/summa* used with the unpartitioned set, semantic responses were marginally faster than pragmatic responses. The weakness of the effect may have been due to data sparseness as there were relatively few pragmatic responders overall. However, the result did not hold up within participants, where more data were available. Moreover, for all set sizes, consistently pragmatic responders were marginally slower to respond to *some/summa* than consistently semantic responders ($\beta = -0.02$, $SE = 0.01$, $p < .09$). Proponents of the Literal-First hypothesis might interpret this as compatible with their claims: By generally taking more time to reach an interpretation of an utterance containing a quantifier, pragmatic responders could be giving themselves the time necessary to generate the scalar inference once they encounter *some* with an unpartitioned set. Some support from this interpretation comes from Bott and Noveck (2004), who showed that participants were more likely to draw a scalar inference when they were given more time to make their judgment. Similarly, De Neys and Schaeken (2007) found that placing participants under increased cognitive load led to fewer scalar implicatures.

However, there is an alternative interpretation. The delay might not be caused by the *computation* of an interpretation, but rather by the time it takes to verify that interpretation in the relevant context. That is, pragmatic responders might be investing more effort in carefully evaluating whether the observed utterance is true in the visual context. We cannot tease apart these two different explanations in our data set, but see Degen and

Tanenhaus (under review) for evidence that semantic and pragmatic responders have different verification strategies.

A further alternative explanation for the pragmatic delay effect may be that giving the pragmatic response is confounded with giving a NO response. If participants are slower to reject a statement than they are to accept it, this would predict NO responses to be slower than YES responses. While it is indeed the case that in our data set NO responses were generally slower than YES responses,⁷ there were multiple quantifiers for which NO responses were faster than YES responses or did not differ in response time (e.g., *none*, *five to twelve*, *some*). Thus, it is unlikely that a response bias was fully responsible for the delay in pragmatic over semantic responses.

5. General discussion

Motivated by the observation that the presence of exact number as an alternative seems to result in delayed responses to the pragmatic interpretation of partitive *some* for small sets, we developed a gumball paradigm to explore the relative naturalness of quantifiers and number terms across set sizes ranging from zero to thirteen. We begin by summarizing the pattern of results across the three experiments. We then discuss the most important theoretical and methodological implications of these results for an understanding of the processing of scalar implicatures.

5.1. Summary of results

In rating studies in which the upper chamber began with 13 gumballs, we found the following results: (1) *some* was less natural for small sets (1–3) than for intermediate size sets (6–8); (2) exact number was more natural than *some*, with the effect being most pronounced for small and large set sizes; (3) descriptions with *some* were less natural than descriptions with *all* for the unpartitioned set, with the effect more pronounced for partitive *some of the* than for simple *some*; (4) intermixing exact number descriptions further lowered naturalness ratings for small sets (one, two, and three gumballs) but not for intermediate size sets.

Result (3) was replicated in a two alternative sentence verification task: the greatest number of YES responses occurred when utterances with *all* were used with the unpartitioned set, fewer with simple *some*, and fewer yet with partitive *some of the*.

Response times in the sentence verification task reflected the patterns from the naturalness rating studies. For YES responses, both response times for participants who on average interpreted *some* semantically (as meaning *some and possibly all*) and response times for participants who interpreted *some* pragmatically increased when *some* was used to refer to the unpartitioned set and to small sets. In addition, response time increased with greater uncertainty about the QUD, which was defined as the consistency with which a participant interpreted *some* semantically or pragmatically. Finally, pragmatic responses to *some* were marginally slower than semantic responses on the subset of the data that included only consistent responders.

5.2. Conclusions

Our results are incompatible with both of the most influential approaches to the processing of scalar implicatures. According to the Literal-First hypothesis, the implicature for the upper-bound interpretation of *some* is computed only after the semantic interpretation of *some*, predicting that NO responses to *some* with the unpartitioned set should always be slower than YES responses. This was indeed the marginally significant result we observed when comparing YES and NO responses between (entirely consistent) participants. Crucially, however, it was not reliable within participants. This cannot be attributed to lack of power. There were three times as many NO judgments in the within-participants analysis compared to the between-participants analysis. Moreover, we cannot rule out the possibility that participants took longer to respond pragmatically than semantically because they are more careful in verifying the utterance in context. The fact that pragmatic responders are generally slower (even in YES responses) than semantic responders is more consistent with a verification hypothesis.

The Literal-First account has trouble explaining the naturalness effects we observed for the YES responses, in particular the slower response time for YES responses to *some/summa* at the unpartitioned set compared to the preferred range. The Literal-First hypothesis predicts that as soon as the first stage of semantic processing is complete and the lower bound verified, a YES response can be made. Thus, according to this account, there should be no difference in YES response times to *some/summa* from 1 to 13 gumballs, where the semantic interpretation holds. Yet we observed clear response time differences.

Proponents of the Literal-First hypothesis might argue that the slowdown effect at the unpartitioned set is due to participants having initially computed the semantic interpretation and then the pragmatic interpretation, before reverting to the semantic interpretation. That is, rather than the semantic response *itself* taking longer to compute, it is the integration with context that might result in an intermediate computation and subsequent cancellation of the implicature that led to the delayed YES response. However, this type of explanation cannot account for the slowdown effects observed for the small sets. For small sets, YES responses to *some/summa* were again slower than responses to number terms. Moreover, the difference in response times was a function of their naturalness relative to number terms. Thus, the more parsimonious explanation is to assume that the effects in both cases arise from the same underlying cause: There is a more natural alternative to *some* that the speaker could have used, but did not.

Our results are also incompatible with the Default model, which assumes that scalar implicatures are computed by default, upon encountering *some*. The Default approach also cannot account for effects of naturalness. Under the Default model, contextual factors like the naturalness of an utterance with a particular quantifier should only come into play when the listener is deciding whether to cancel the already computed implicature. That is, YES responses to the unpartitioned set should be the only set size where naturalness of *some* affects response times. To see this, consider what interpretation is necessary for each set size to arrive at a YES response. For the unpartitioned set, only the semantic *some and possibly all* interpretation yields a YES response. In contrast, for set sizes

1–12, both the semantic and the pragmatic interpretation yield YES responses. Thus, under the Default model YES responses to *some* should have been equally long for set sizes 1–12, and a slowdown is expected only for the unpartitioned set, where an additional cost for canceling the implicature is incurred.

The Default account correctly predicts the slowdown effect for YES responses to *some* with the unpartitioned set compared to its most preferred range. However, it cannot explain (a) the slowdown effect for small sets and (b) the fact that pragmatic NO responses to the unpartitioned set are also slower compared to YES responses in the preferred range.

We thus believe that the pattern of results we obtained in Experiments 1–3 is most compatible with a Constraint-Based account in which the speed and robustness of an implicature is determined by the probabilistic support it receives from multiple cues available in the linguistic and discourse context, including the task/goal relevant information. In this paper, we investigated the effect of alternatives on processing *some* and the scalar inference from *some* to *not all*. In particular, we investigated (a) the syntactic partitive as a cue to the implicature and (b) the naturalness and availability of lexical alternatives to *some* as inhibitory cues to arriving at an interpretation for *some*.

Some observations on the naturalness/availability of lexical alternatives are in order. While linguists and logicians treat the meaning of *some*, *all*, and other quantifiers as set-theoretically well-defined as for example in Generalized Quantifier Theory (Barwise & Cooper, 1981), the naturalness results from Experiments 1 and 2 show that not only are quantifiers like *some* more natural for some set sizes than others, but that their naturalness depends crucially on alternative lexical items that the speaker might have used. That is, while there may be ways to unambiguously define quantifier meanings, quantifier interpretation is a matter of degree. Some set sizes (e.g., 6 or 7 out of 13 gumballs) are better fits for *some* than others (e.g., 2 or 13 out of 13 gumballs) and introducing alternatives can change the goodness of fit. This is paralleled in the concepts and categories literature, where typicality influences categorization even for logically well-defined concepts such as parity (Armstrong et al., 1983). Some numbers are judged as better instances of an odd or an even number than others. Similarly, in our studies *some* is judged as a more or less natural label for different sizes (and processed more or less quickly accordingly), despite its logical definition unambiguously assigning a TRUE or FALSE value for any given set size.

The effects of naturalness and availability are also compatible with accounts that treat the meaning of quantifiers as distributions over quantities. These distributions reflect listener beliefs about the probability of a particular quantifier being uttered, given a particular set size. Listener beliefs are updated by different contextual factors, among them the availability of alternatives. That is, the distributions may shift contextually depending on the available alternatives. Mapping this onto our naturalness results for small sets: the posterior probability of observing an utterance of *some* to refer to a set of size 2 given that number terms are contextually available is lower than when they are not.

Methodologically, this means that it is important for researchers to be aware of the relative naturalness of the quantifiers under investigation, that is, which range of the interpretive

distribution one is sampling from. For example, in light of the intrusion effects we have shown, delays in the processing of *some* observed by Huang and Snedeker (2009) could have an explanation that is unrelated to the processing of *implicatures* per se: Not only did they use a set size that is relatively unnatural for *some* for (two), but *two* was also explicitly included as a lexical alternative among the experimental items. Therefore, *some* may have been a dispreferred label for referring to the set size it was used for and thus caused the delay. Note that Huang and Snedeker did not find a similar delay for *all*,⁸ which was used with a larger set size (3). Based on our finding, we would predict a delay for *all* if it had been paired with a set size of 2 and a smaller delay for *some* if it had been paired with a larger set size. These predictions are confirmed in Degen and Tanenhaus (unpublished data).

Another important factor that is likely to affect the speed and robustness of scalar implicatures is the QUD. While the QUD was not explicitly manipulated in our studies, it may have played a substantial role in participants' response behavior. In Experiment 3, we interpreted participants' response consistency to *some* used with the unpartitioned set as indicative of how much uncertainty they had about the QUD. Our reasoning was that an implicature from *You got some of the gumballs* to *You got some, but not all of the gumballs* should be more likely if the assumed QUD is *Did I get all of the gumballs?* rather than *Did I get any gumballs?* (see Zondervan, 2008 for evidence that both implicit and explicit QUDs affect implicature rates). Thus, we took within-participant distributions over semantic and pragmatic responses to reflect the uncertainty that participants had about the QUD—more skewed distributions, whether skewed toward semantic or pragmatic responses, were taken to indicate less uncertainty. We found that participants with less uncertainty were faster to respond, regardless of whether they responded semantically or pragmatically. This provides an interesting further testing ground for the Constraint-Based account, which predicts that given a QUD that makes the stronger alternative *You got all of the gumballs* particularly relevant, the implicature should be more rapidly available than in the neutral case (as in our Experiment 3) and in turn the implicature should be even less rapidly available in a lower-bound context. We are currently exploring this hypothesis. More generally, our results suggest that it will be important for future research to explicitly manipulate the QUD.

While in this paper we have focused on the naturalness and availability of *some* and its alternatives (partitive vs. non-partitive, *some* vs. other quantifiers), and to some extent on the QUD, there are many other cues that listeners are likely to take into consideration when making an interpretive choice about an utterance with *some*, some of which have been previously shown to affect the rate and processing speed of scalar implicatures. Among them are cognitive load (De Neys & Schaeken, 2007), threat of loss of face (Bonnefon, Feeney, & Villejoubert, 2009), speaker competence and reliability (Grodner & Sedivy, 2011), average informativeness of the stronger alternative (Stiller, Goodman, & Frank, 2011), semantic context (Chierchia, 2004), and prosody (de Marneffe & Tonhauser, 2014).

Taken together, then, the evidence suggests that listeners, upon encountering a scalar item, simultaneously consider information from multiple cues in the linguistic and discourse context in determining whether an implicature is likely, and if so, what form it should take in a particular context. Rather than granting privileged status to the semantic interpretation

of *some*, under a Constraint-Based account the research program becomes one of identifying and quantifying the cues that listeners use in computing scalar implicatures. The goal, then, is to provide a unified account of why in many situations this computation appears to be delayed and to come at a cost, whereas in others situations, it is more rapid, and less resource-intensive. The work reported here is a first step in that direction.

We conclude with two observations, one theoretical and one methodological. The theoretical observation is that the perspective we are pursuing presents a challenge to the classic distinction between Generalized and Particularized Conversational Implicatures. If we are correct about the context dependency of GCIs, then they no longer have a special status relative to other conversational implicatures. This is analogous to related claims in the literature, for example, the view that lexical representations are distributed and processing a word always involves a context-dependent computation, rather than retrieving a “core” meaning (Elman, 2004; MacDonald & Seidenberg, 2006).

The methodological observation concerns the relationship between formal theories and research in Experimental Pragmatics. We strongly endorse the view that Experimental Pragmatics must be informed by formal theories. However, research in Experimental Pragmatics requires grounding the experimental work in a task, which necessarily involves mapping utterances onto specific contexts. Experimental tests of those theories will often result in novel observations and new data that are central to understanding pragmatics, but might be orthogonal to our current theories. We view this as part of the challenge and excitement of current work in the field.

Indeed, the effects of numbers statements on upper-bound *some* is an example of a specific result that arises in particular experimental contexts. However, it has more general implications for the importance of the availability and naturalness of alternatives in guiding implicature and, more generally, for the hypothesis that Generalized Implicatures differ in kind from Particularized Implicatures.

Acknowledgments

This work has benefited from conversations with Dan Grodner, Florian Jaeger, Richard Breheny, Jesse Snedeker, Yi Ting Huang, and members of the Tanenhaus lab. We thank Dana Subik for testing participants and Patricia Reeder for being a competent voice. Special thanks to our action editor, Delphine Dahan, for cogent comments that greatly improved the manuscript, as well as two anonymous reviewers. This work was partially supported by NIH grant HD 27206.

Notes

1. This holds only *ceteris paribus*. Chierchia (2004) and others, for example, have noted that in downward-entailing contexts, where monotonicity properties are

reversed, scalar implicatures (as in 3a) do not arise. In addition, a scale reversal takes place, often resulting in weaker lower-bound implicatures (as in 3b):

(3) It is not the case that Joe ate most of the cookies.

(a) \hookrightarrow It is not the case that Joe ate most but not all of the cookies

= Joe ate all of the cookies.

(b) Joe ate at least some of the cookies.

2. However, see the discussion surrounding the relevance of the stronger scalar alternative in section 4.3.
3. In Fig. 1, further potential cues that listeners may make use of are listed. The QUD (Roberts, 1996) will be discussed further in section 4.3. Prosody may affect the interpretation because contrastive stress on *some* forces the implicature. In contrast, monotonicity properties of the context have been shown to affect scalar implicature computation: Scalar implicatures should not arise (or arise much more weakly) in downward-entailing contexts, as discussed above.
4. To test whether the naturalness difference between *some/summa* used with small sets versus used with a set size of 6, we performed an additional analysis whereby naturalness differences were generated by subtracting the small set naturalness ratings from ratings that were randomly sampled from the distribution of ratings on set size 6 trials. These differences were then predicted from set size. As set size increased, the differences decreased ($\beta = -0.76$, $SE = 0.10$, $\chi^2(1) = 44.94$, $p < .0001$).
5. Both Noveck and Posada (2003) and Bott and Noveck (2004) found that participants tended to respond YES or NO consistently to underinformative items. For example in Noveck and Posada (2003)'s study, there was one group of 7 consistently semantic responders (37%) and another group of 12 consistently pragmatic responders (63%). Similarly, in Bott and Noveck (2004)'s Study 3, 41% of responses to underinformative items were semantic while 59% were pragmatic. However, participants within these categories were not always entirely consistent. Noveck and Posada note that their semantic responders were 96% consistent in their answers, while their pragmatic responders were 92% consistent. That is, some responders classified as semantic also gave some pragmatic responses and vice versa. Similarly, Bott and Noveck found that of their 32 participants, only nine participants gave entirely consistent answers (two semantic, seven pragmatic). The remaining 23 participants gave both types of answers.
6. See Appendix B for an analysis of participants' correlated response behavior at the upper bound and their propensity to draw plural implicatures, that is, reject the *some* statement when they received one gumball.
7. This was assessed with a mixed effects linear regression predicting log-transformed response times in the entire data set from response type (YES or NO, $\beta = -0.02$, $SE = 0.01$, $t = -2.4$).

8. Though they do report a delay for *all* relative to *three* in their Experiment 3, which they explain by reference to the increased difficulty involved in appropriate domain restriction in this experiment.

References

- Armstrong, S. L., Gleitman, L. R., & Gleitman, H. (1983). What some concepts might not be. *Cognition*, 13 (3), 263–308.
- Atkinson, J., Campbell, F. W., & Francis, M. R. (1976). The magic number 4 ± 0 : A new look at visual numerosity judgements. *Perception*, 5(3), 327–334.
- Baayen, R. H., Davidson, D. J., & Bates, D. M. (2008). Mixed-effects modeling with crossed random effects for subjects and items. *Journal of Memory and Language*, 59(4), 390–412. doi:10.1016/j.jml.2007.12.005.
- Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013). Random effects structure in mixed-effects models: Keep it maximal. *Journal of Memory and Language*, 68(3), 255–278.
- Barwise, J., & Cooper, R. (1981). Generalized quantifiers and natural language. *Linguistics and Philosophy*, 4 (2), 159–219.
- Bonnefon, J.-F., Feeney, A., & Villejoubert, G. (2009). When some is actually all: Scalar inferences in face-threatening contexts. *Cognition*, 112(2), 249–258. doi:10.1016/j.cognition.2009.05.005.
- Bott, L., Bailey, T. M., & Grodner, D. (2012). Distinguishing speed from accuracy in scalar implicatures. *Journal of Memory and Language*, 66(1), 123–142. doi:10.1016/j.jml.2011.09.005.
- Bott, L., & Noveck, I. (2004). Some utterances are underinformative: The onset and time course of scalar inferences. *Journal of Memory and Language*, 51(3), 437–457. doi:10.1016/j.jml.2004.05.006.
- Bradlow, A. R., & Bent, T. (2008). Perceptual adaptation to non-native speech. *Cognition*, 106(2), 707–729.
- Breheny, R., Katsos, N., & Williams, J. (2006). Are generalised scalar implicatures generated by default? An on-line investigation into the role of context in generating pragmatic inferences. *Cognition*, 100(3), 434–463. doi:10.1016/j.cognition.2005.07.003.
- Carston, R. (1998). Informativeness, relevance and scalar implicature. In R. Carston & S. Uchida (Eds.), *Relevance theory: Applications and implications* (pp. 179–236). Amsterdam: John Benjamins.
- Chambers, C. G., Tanenhaus, M. K., Eberhard, K. M., Filip, H., & Carlson, G. N. (2002). Circumscribing referential domains during real-time language comprehension. *Journal of Memory and Language*, 47(1), 30–49. doi:10.1006/jmla.2001.2832.
- Chambers, C. G., Tanenhaus, M. K., & Magnuson, J. S. (2004). Actions and affordances in syntactic ambiguity resolution. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 30(3), 687–696.
- Chang, F., Dell, G. S., & Bock, K. (2006). Becoming syntactic. *Psychological Review*, 113(2), 234–272. doi:10.1037/0033-295X.113.2.234.
- Chierchia, G. (2004). Scalar implicatures, polarity phenomena, and the syntax/pragmatics interface. In A. Belletti (Ed.), *Structures and beyond* (Vol. 3, pp. 39–103). Oxford, England: Oxford University Press.
- Clark, A. (2013). Whatever next? Predictive brains, situated agents, and the future of cognitive science. *Behavioral and Brain Sciences*, 36(3), 181–204.
- Clayards, M., Tanenhaus, M. K., Aslin, R. N., & Jacobs, R. A. (2008). Perception of speech reflects optimal use of probabilistic speech cues. *Cognition*, 108(3), 804–809. doi:10.1016/j.cognition.2008.04.004.
- Cooper, R. M. (1974). The control of eye fixation by the meaning of spoken language: A new methodology for the real-time investigation of speech perception, memory, and language processing. *Cognitive Psychology*, 6(1), 84–107.
- de Marneffe, M., & Tonhauser, J. (2014). Prosody affects scalar implicature generation. Poster, 27th CUNY Conference on Human Sentence Processing, Columbus, OH.

- De Neys, W., & Schaeken, W. (2007). When people are more logical under cognitive load – Dual task impact on scalar implicature. *Experimental Psychology*, 54(2), 128–133. doi:10.1027/1618-3169.54.2.128.
- Degen, J., & Tanenhaus, M. K. (under review). Availability of alternatives and the processing of scalar implicatures: a visual world eye-tracking study.
- Elman, J. L. (2004). An alternative view of the mental lexicon. *Trends in Cognitive Sciences*, 8(7), 301–306. doi:10.1016/j.tics.2004.05.003.
- Fine, A. B., Jaeger, T. F., Farmer, T. F., & Qian, T. (2013). Rapid expectation adaptation during syntactic comprehension. *PLoS ONE*, 8(10), e77661, doi:10.1371/journal.pone.0077661
- Gazdar, G. (1979). *Pragmatics: Implicature, presupposition, and logical form*. New York: Academic Press.
- Green, M. S. (1995). Quantity, volubility, and some varieties of discourse. *Linguistics and Philosophy*, 18(1), 83–112.
- Grice, H. P. (1975). Logic and conversation. *Syntax and Semantics*, 3, 41–58.
- Grodner, D. J., Klein, N. M., Carberry, K. M., & Tanenhaus, M. K. (2010). “Some,” and possibly all, scalar inferences are not delayed: Evidence for immediate pragmatic enrichment. *Cognition*, 116(1), 42–55. doi:10.1016/j.cognition.2010.03.014.
- Grodner, D. J., & Sedivy, J. C. (2011). The effect of speaker-specific information on pragmatic inferences. In N. Pearlmuter & E. Gibson (Eds.), *The processing and acquisition of reference* (Vol. 2327, pp. 239–272). Cambridge, MA: MIT Press.
- Horn, L. (1984). Toward a new taxonomy for pragmatic inference: Q-based and R-based implicature. In D. Schiffrin (Ed.), *Meaning, form, and use in context: Linguistic applications* (pp. 11–42). Washington, DC: Georgetown University Press.
- Horn, L. (2004). Implicature. In L. Horn & G. Ward (Eds.), *Handbook of pragmatics* (pp. 3–28). Oxford, England: Blackwell.
- Huang, Y. T., Hahn, N., & Snedeker, J. (2010). Some inferences still take time: Prosody, predictability, and the speed of scalar implicatures. *Poster Presented at the 23rd Annual CUNY Conference on Human Sentence Processing*.
- Huang, Y. T., & Snedeker, J. (2009). Online interpretation of scalar quantifiers: Insight into the semantics-pragmatics interface. *Cognitive Psychology*, 58(3), 376–415. doi:10.1016/j.cogpsych.2008.09.001.
- Huang, Y. T., & Snedeker, J. (2011). Logic and conversation revisited: Evidence for a division between semantic and pragmatic content in real-time language comprehension. *Language and Cognitive Processes*, 26(8), 1161–1172.
- Kahneman, D. (2011). *Thinking, fast and slow*. New York: Farrar, Straus, and Giroux.
- Kaufman, E. L., Lord, M. W., Reese, T. W., & Volkman, J. (1949). The discrimination of visual number. *The American Journal of Psychology*, 62(4), 498–525.
- Kurumada, C., Brown, M., & Tanenhaus, M. K. (2012). Pragmatic interpretation of contrastive prosody: It looks like speech adaptation. In N. Miyake, D. Peebles, & R. P. Cooper (Eds.), *Proceedings of the 34th Annual Conference of the Cognitive Science Society* (pp. 647–652). Austin, TX: Cognitive Science Society.
- Kutas, M., & Hillyard, S. A. (1980). Reading senseless sentences: Brain potentials reflect semantic incongruity. *Science*, 207(4427), 203–205.
- Levinson, S. (2000). *Presumptive meanings*. Cambridge, England: Cambridge University Press.
- Levy, R. (2008). Expectation-based syntactic comprehension. *Cognition*, 106(3), 1126–1177. doi:10.1016/j.cognition.2007.05.006.
- MacDonald, M., Pearlmutter, N., & Seidenberg, M. (1994). The lexical nature of syntactic ambiguity resolution. *Psychological Review*, 101, 676–703.
- MacDonald, M. C., & Seidenberg, M. S. (2006). Constraint satisfaction approaches to lexical and sentence comprehension. In M. A. Gernsbacher & M. J. Traxler (Eds.), *Handbook of psycholinguistics* (pp. 581–612). London: Elsevier.
- Mandler, G., Shebo, B. J., & Vol, I. (1982). Subitizing: An analysis of its component processes. *Journal of Experimental Psychology: General*, 111(1), 1–22.

- Neely, J. H. (1977). Semantic priming and retrieval from lexical memory: Roles of inhibitionless spreading activation and limited-capacity attention. *Journal of Experimental Psychology: General*, 106(3), 226–254. doi:10.1037//0096-3445.106.3.226.
- Noveck, I., & Posada, A. (2003). Characterizing the time course of an implicature: An evoked potentials study. *Brain and Language*, 85(2), 203–210. doi:10.1016/S0093-934X(03)00053-1.
- Posner, M. I., & Snyder, C. R. (1975). Facilitation and inhibition in the processing of signals. In P. M. A. Rabbitt, & S. Dornic (Eds.), *Attention and performance* (pp. 669–682). New York: Academic Press.
- Roberts, C. (1996). Information structure in discourse: Towards an integrated formal theory of pragmatics. In J. H. Yoon & A. Kathol (Eds.), *OSU working papers in linguistics 49: Papers in semantics* (pp. 91–136). Columbus: The Ohio State University.
- Russell, B. (1905). On denoting. *Mind*, 14(4), 479–493. doi:10.1093/mind/XIV.4.479.
- Russell, B. (2012). *Probabilistic reasoning and the computation of scalar implicatures*. Doctoral dissertation: Brown University.
- Sauerland, U. (2004). Scalar implicatures in complex sentences. *Linguistics and Philosophy*, 27(3), 367–391. doi:10.1023/B:LING.0000023378.71748.db.
- Sedivy, J. C., Tanenhaus, M. K., Chambers, C. G., & Carlson, G. N. (1999). Achieving incremental semantic interpretation through contextual representation. *Cognition*, 71(2), 109–147.
- Seidenberg, M., & Macdonald, M. C. (1999). A probabilistic constraints approach to language acquisition and processing. *Cognitive Science*, 23(4), 569–588.
- Shiffrin, R. M., & Schneider, W. (1977). Controlled and automatic human information processing. *Psychological Review*, 84, 127–190.
- Sperber, D., & Wilson, D. (1995). *Relevance: Communication and cognition*. Oxford, England: Blackwell.
- Stiller, A., Goodman, N. D., & Frank, M. C. (2011). Ad-hoc scalar implicature in adults and children. In L. Carlson, C. Holscher & T. Shipley (Eds.), *Proceedings of the 33rd Annual Meeting of the Cognitive Science Society*. Austin, TX: Cognitive Science Society.
- Tanenhaus, M. K., Spivey-Knowlton, M. J., Eberhard, K. M., & Sedivy, J. C. (1995). Integration of visual and linguistic information in spoken language comprehension. *Science*, 268, 1632–1634.
- Trueswell, J. C., Tanenhaus, M. K., & Kello, C. (1993). Verb-specific constraints in sentence processing: Separating effects of lexical preference from garden-paths. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 19 (3), 528–553. Available at: <http://www.ncbi.nlm.nih.gov/pubmed/8501429>.
- Tanenhaus, M.K., & Trueswell, J. C. (1995). Sentence comprehension. In J. Miller & P. Elmas (Eds.), *Speech, language, and communication* San Diego, CA: Academic Press.
- Van Rooij, R., & Schulz, K. (2004). Exhaustive interpretation of complex sentences. *Journal of Logic, Language and Information*, 13(4), 491–519.
- Zondervan, A. (2009). Experiments on QUD and focus as a contextual constraint on scalar implicature calculation. In U. Sauerland, & K. Yatsushiro (Eds.), *Semantics and Pragmatics: From Experiment to Theory*, Ch. 6. Basingstoke: Palgrave-Macmillan.
- Zweig, E. (2009). Number-neutral bare plurals and the multiplicity implicature. *Linguistics and Philosophy*, 32(4), 353–407.

Appendix A

Table A1

Set sizes sampled by the 12 base lists used in Experiments 1 and 2

List	Set Sizes	List	Set Sizes
1	1, 2, 5, 9	7	2, 3, 5, 9
2	1, 2, 6, 10	8	2, 3, 6, 10
3	1, 3, 5, 9	9	2, 4, 7, 11
4	1, 3, 7, 11	10	2, 4, 8, 12
5	1, 4, 6, 10	11	3, 4, 7, 11
6	1, 4, 8, 12	12	3, 4, 8, 12

Appendix B

We observed an interesting correlation between participants' response behavior to *some/summa* used with the unpartitioned set (upper bound) versus when it was used for one gumball (lower bound). At the lower bound, the sentences *You got some (of the) gumballs* are strictly speaking true but often trigger a multiplicity implicature, whereby the cardinality of the set denoted by *some (of the) gumballs* is expected to be greater than one (Zweig, 2009). We observe response behavior to *some* at the lower bound similar to that at the upper bound: 52% of responses were YES responses, while 48% were NO responses. Of these, 38 participants (81%) were completely consistent, 8 participants (17%) gave one response that was different from the rest, and only one participant gave




Fig. B1. Jittered scatterplot of participants' response behavior at the upper and lower bound. Each point represents one participant. Horizontal line represents the perfect correlation line.

50% pragmatic and 50% semantic responses at the lower bound. Interestingly, 88% of participants who responded pragmatically at the upper bound also responded pragmatically at the lower bound. In addition, 40% of participants who had responded semantically or inconsistently at the upper bound responded pragmatically at the lower bound. That is, most participants who responded pragmatically at the upper bound also responded pragmatically at the lower bound, and most people who responded semantically at the lower bound also responded semantically at the upper bound (see Fig. B1). The correlation between the number of semantic responses a participant gave at the upper bound versus at the lower bound is 0.45.